

TASK #1: UNDERSTAND THE PROBLEM STATEMENT AND BUSINESS CASE

- In this project, we will predict real estate prices using artificial neural networks.
- Dataset includes house sale prices for King County in USA and homes that are sold in the time period between May, 2014 and May, 2015.
- Columns:
 - `ida`: notation for a house
 - `date`: Date house was sold
 - `price`: Price is prediction target
 - `bedrooms`: Number of Bedrooms/House
 - `bathrooms`: Number of bathrooms/House
 - `sqft_living`: square footage of the home
 - `qft_lot`: square footage of the lot
 - `floors`: Total floors (levels) in house
 - `waterfront`: House which has a view to a waterfront
 - `view`: Has been viewed



- Data Source: <https://www.kaggle.com/harlfoxem/housesalesprediction>
- Image Source: <https://pxhere.com/en/photo/1440167>

- Columns:
 - `condition`: How good the condition is
 - `grade`: overall grade given to the housing unit, based on King County grading system
 - `sqft_abovesquare`: footage of house apart from basement
 - `sqft_basement`: square footage of the basement
 - `yr_built`: Built Year
 - `yr_renovated`: Year when house was renovated
 - `zipcode`: zip
 - `lat`: Latitude coordinate
 - `long`: Longitude coordinate
 - `sqft_living15`: Living room area in 2015(implies some renovations)
 - `sqft_lot15`: lot Size area in 2015 (implies some renovations)

- Data Source:
<https://www.kaggle.com/harlfoxem/housesalesprediction>

INSTRUCTOR

- Adjunct professor & online instructor
- Passionate about artificial intelligence, machine learning, and electric vehicles
- Taught 250,000+ students globally
- MBA (2018), Ph.D. (2014), M.A.Sc (2011)



Ryan Ahmed, Ph.D.

TASK #2: IMPORT LIBRARIES AND DATASETS

```
In [3]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# from jupyterthemes import jtplot
# jtplot.style(theme='monokai', context='notebook', ticks=True, grid=False)
# # setting the style of the notebook to be monokai theme
# # this line of code is important to ensure that we are able to see the x and y axes clearly
# # If you don't run this code line, you will notice that the xlabel and ylabel on any plot is black on black and it
```

```
In [4]: house_df = pd.read_csv('realestate_prices.csv', encoding = 'ISO-8859-1')
```

```
In [5]: house_df
```

```
Out[5]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	yr
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650	1.0	0	0	...	7	1180	0	0
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242	2.0	0	0	...	7	2170	400	0
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000	1.0	0	0	...	6	770	0	0
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000	1.0	0	0	...	7	1050	910	0
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080	1.0	0	0	...	8	1680	0	0
...
21608	263000018	20140521T000000	360000.0	3	2.50	1530	1131	3.0	0	0	...	8	1530	0	0
21609	6600060120	20150223T000000	400000.0	4	2.50	2310	5813	2.0	0	0	...	8	2310	0	0
21610	1523300141	20140623T000000	402101.0	2	0.75	1020	1350	2.0	0	0	...	7	1020	0	0
21611	291310100	20150116T000000	400000.0	3	2.50	1600	2388	2.0	0	0	...	8	1600	0	0

21611	291310100	20150116T000000	400000.0	3	2.50	1600	2388	2.0	0	0	...	8	1600	0
21612	1523300157	20141015T000000	325000.0	2	0.75	1020	1076	2.0	0	0	...	7	1020	0

21613 rows × 21 columns

In [6]: house_df.head(5)

Out[6]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	yr_built
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650	1.0	0	0	...	7	1180	0	1955
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242	2.0	0	0	...	7	2170	400	1951
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000	1.0	0	0	...	6	770	0	1933
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000	1.0	0	0	...	7	1050	910	1968
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080	1.0	0	0	...	8	1680	0	1987

5 rows × 21 columns

In [7]: house_df.tail(10)

Out[7]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	y
21603	7852140040	20140825T000000	507250.0	3	2.50	2270	5536	2.0	0	0	...	8	2270	0	0
21604	9834201367	20150126T000000	429000.0	3	2.00	1490	1126	3.0	0	0	...	8	1490	0	0
21605	3448900210	20141014T000000	610685.0	4	2.50	2520	6023	2.0	0	0	...	9	2520	0	0
21606	7936000429	20150326T000000	1007500.0	4	3.50	3510	7200	2.0	0	0	...	9	2600	910	0
21607	2997800021	20150219T000000	475000.0	3	2.50	1310	1294	2.0	0	0	...	8	1180	130	0
21608	263000018	20140521T000000	360000.0	3	2.50	1530	1131	3.0	0	0	...	8	1530	0	0
21609	6600060120	20150223T000000	400000.0	4	2.50	2310	5813	2.0	0	0	...	8	2310	0	0
21610	1523300141	20140623T000000	402101.0	2	0.75	1020	1350	2.0	0	0	...	7	1020	0	0
21611	291310100	20150116T000000	400000.0	3	2.50	1600	2388	2.0	0	0	...	8	1600	0	0
21612	1523300157	20141015T000000	325000.0	2	0.75	1020	1076	2.0	0	0	...	7	1020	0	0

10 rows × 21 columns

In [8]: house_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               21613 non-null   int64  
 1   date             21613 non-null   object  
 2   price            21613 non-null   float64 
 3   bedrooms         21613 non-null   int64  
 4   bathrooms        21613 non-null   float64 
 5   sqft_living      21613 non-null   int64  
 6   sqft_lot          21613 non-null   int64  
 7   floors            21613 non-null   float64 
 8   waterfront        21613 non-null   int64  
 9   view              21613 non-null   int64  
 10  condition         21613 non-null   int64  
 11  grade             21613 non-null   int64  
 12  sqft_above        21613 non-null   int64  
 13  sqft_basement     21613 non-null   int64  
 14  yr_built          21613 non-null   int64  
 15  yr_renovated      21613 non-null   int64  
 16  zipcode            21613 non-null   int64  
 17  lat                21613 non-null   float64 
 18  long               21613 non-null   float64 
 19  sqft_living15      21613 non-null   int64  
 20  sqft_lot15         21613 non-null   int64 
```

```
20 sqft_lot15    21613 non-null int64
dtypes: float64(5), int64(15), object(1)
memory usage: 3.5+ MB
```

PRACTICE OPPORTUNITY #1 [OPTIONAL]:

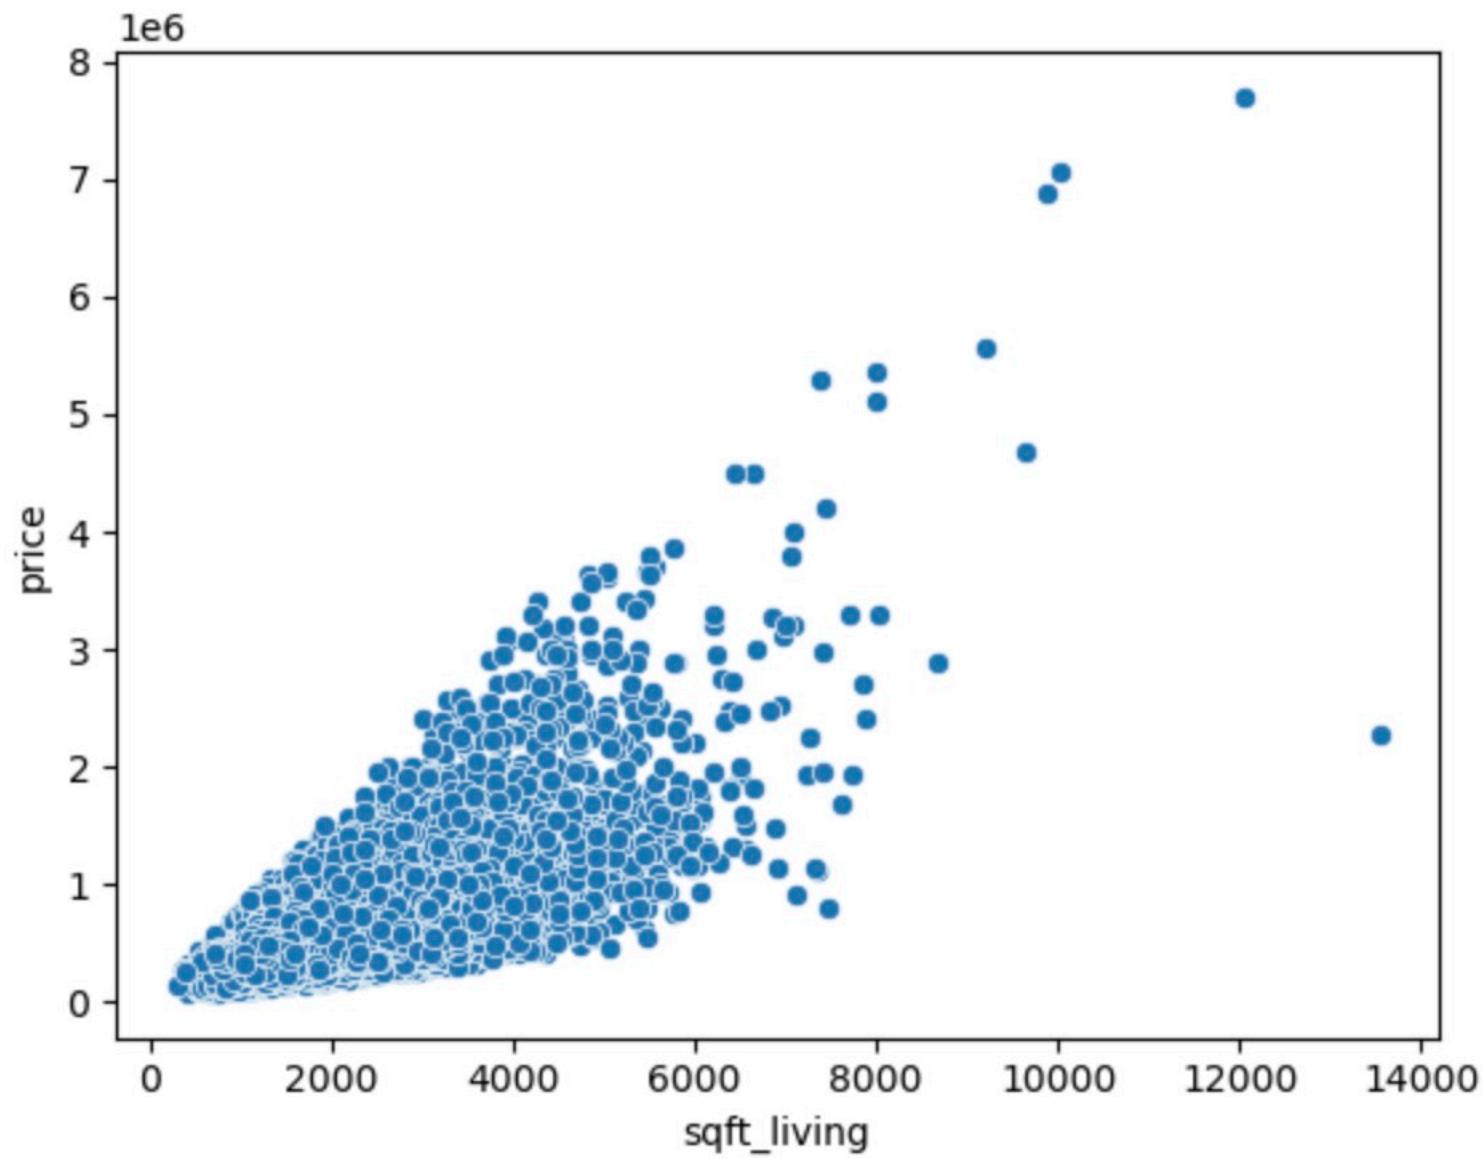
- What is the average house price?
- What is the price of the cheapest house?
- What is the average number of bathrooms and bedrooms? round your answer to the lowest value
- What is the maximum number of bedrooms?

In []:

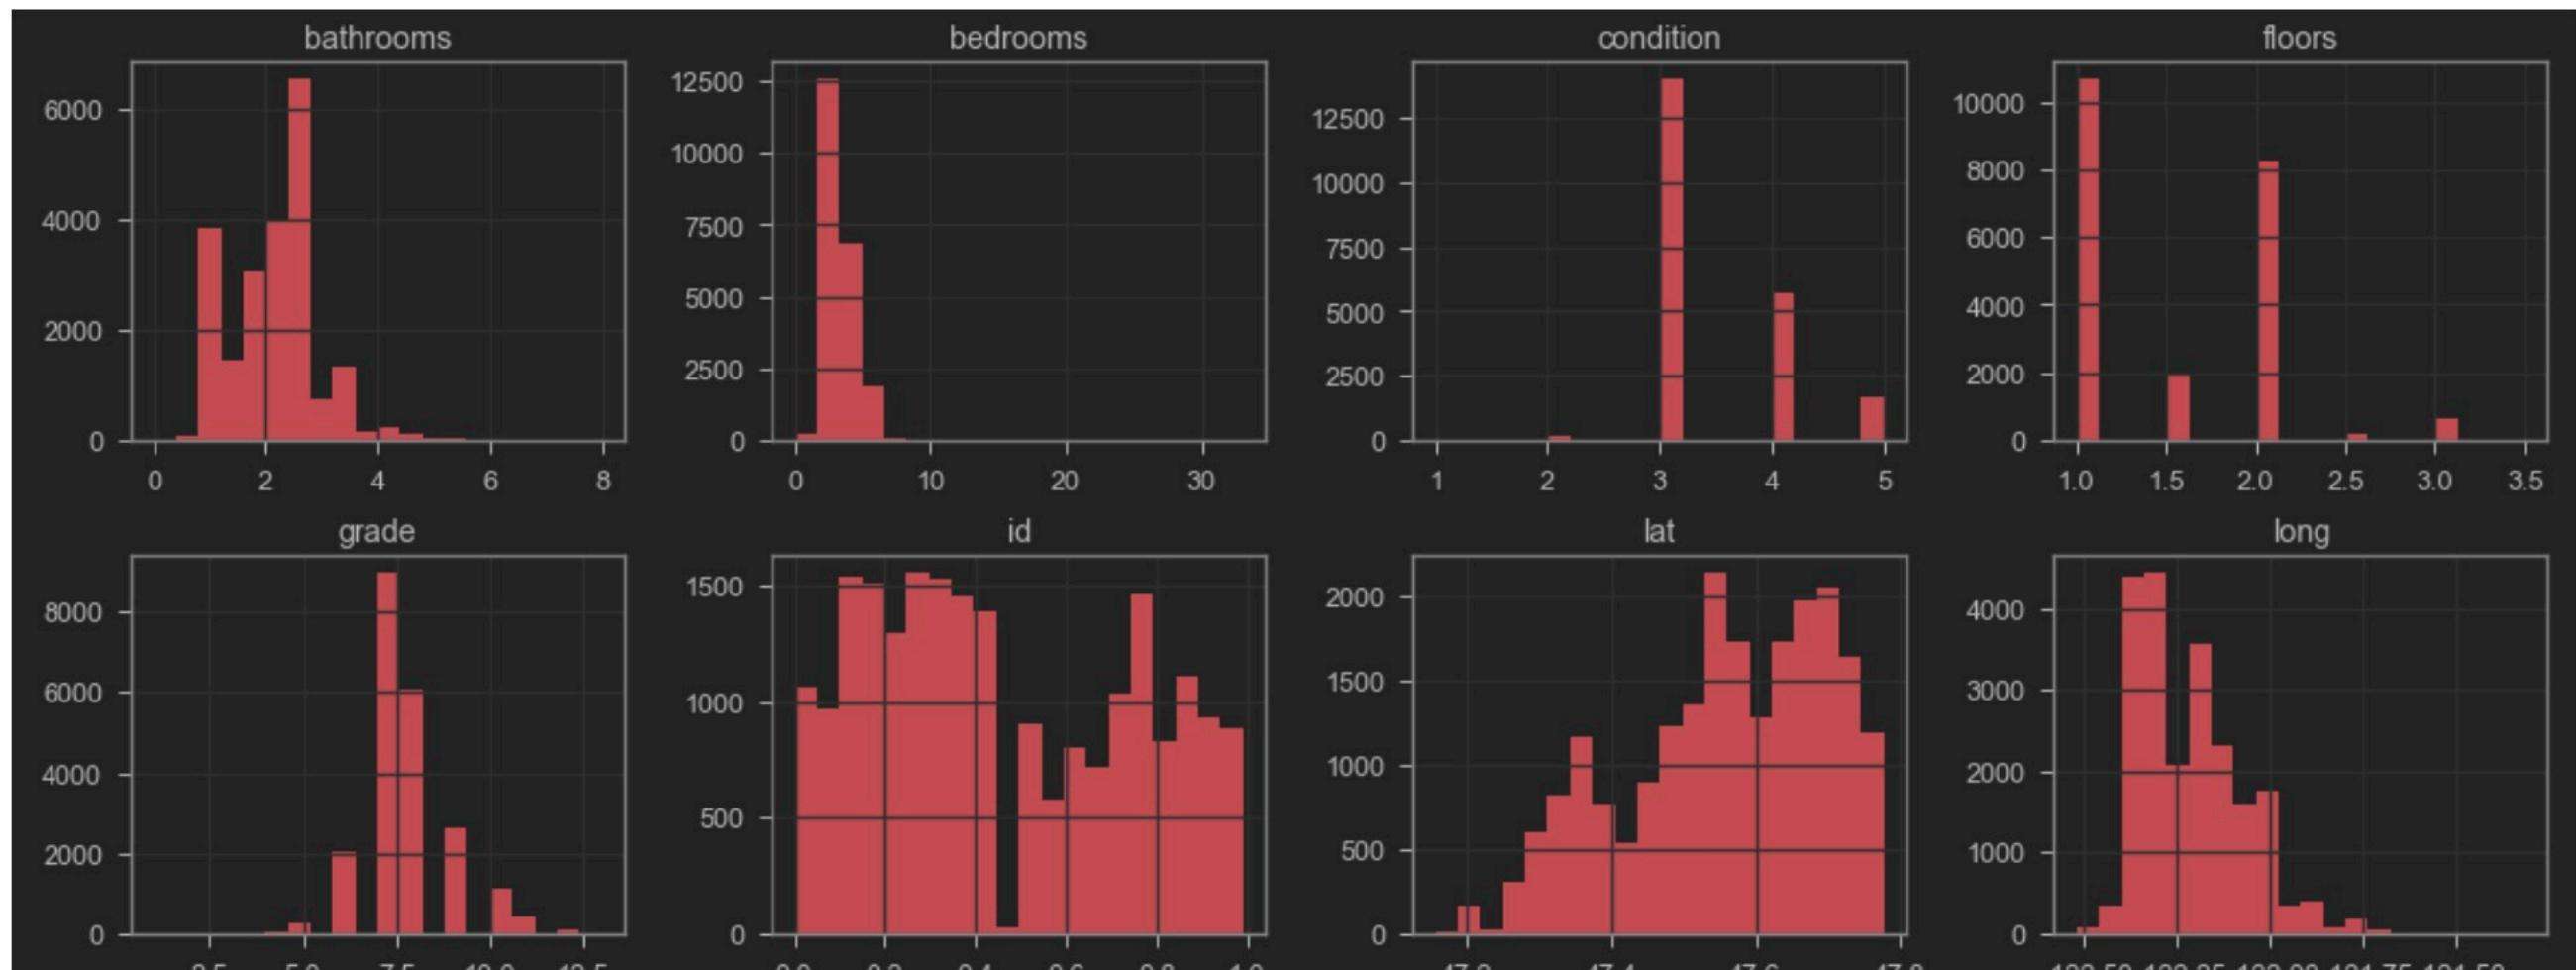
TASK #3: PERFORM DATA VISUALIZATION

```
In [9]: sns.scatterplot(x = 'sqft_living', y = 'price', data = house_df)
```

```
Out[9]: <Axes: xlabel='sqft_living', ylabel='price'>
```



```
In [8]: house_df.hist(bins = 20, figsize = (20,20), color = 'r');
```

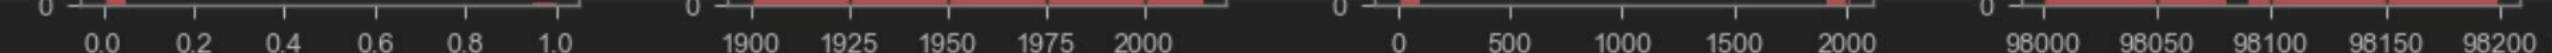




```
In [9]: f, ax = plt.subplots(figsize = (20, 20))
sns.heatmap(house_df.corr(), annot = True)
```

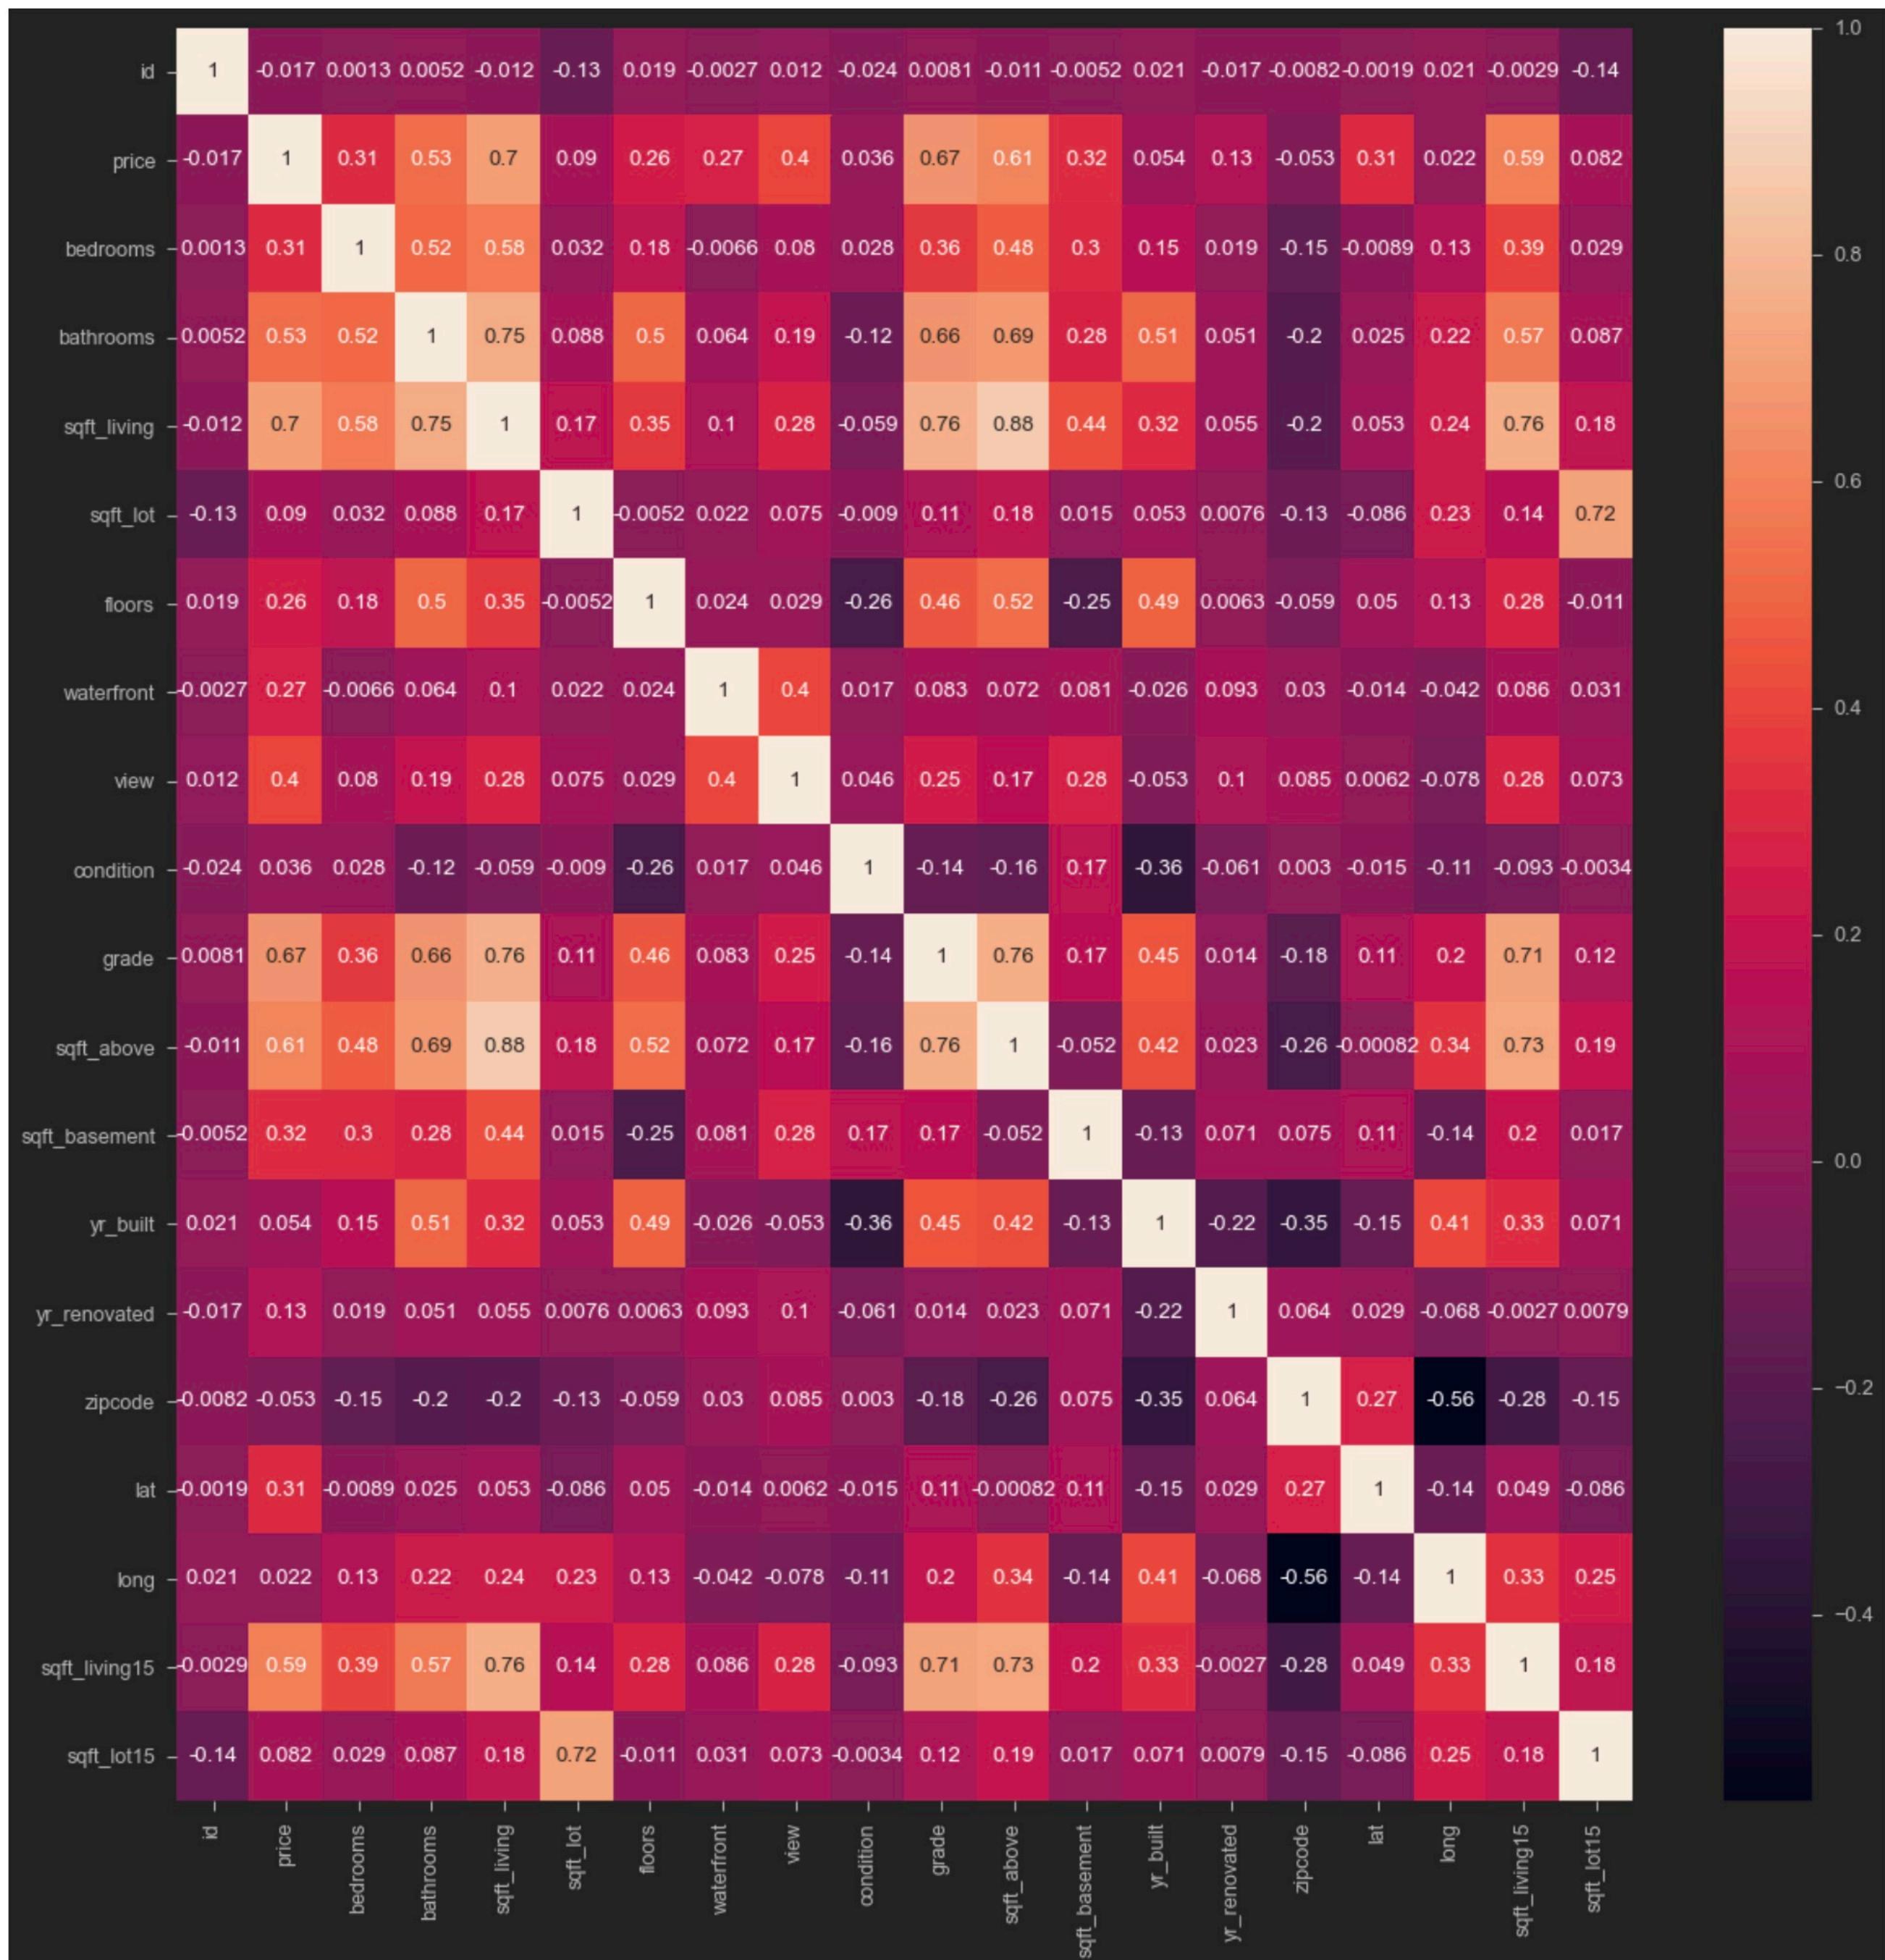
```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x20f60957828>
```





```
In [9]: f, ax = plt.subplots(figsize = (20, 20))
sns.heatmap(house_df.corr(), annot = True)
```

```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x20f60957828>
```



```
In [10]: house_df_sample = house_df[ ['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'sqft_above', 'sqft_baseme
```

```
In [11]: house_df_sample
```

```
In [11]: house_df_sample
```

```
Out[11]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	sqft_above	sqft_basement	yr_built
0	221900.0	3	1.00	1180	5650	1180	0	1955
1	538000.0	3	2.25	2570	7242	2170	400	1951
2	180000.0	2	1.00	770	10000	770	0	1933
3	604000.0	4	3.00	1960	5000	1050	910	1965
4	510000.0	3	2.00	1680	8080	1680	0	1987
...
21608	360000.0	3	2.50	1530	1131	1530	0	2009
21609	400000.0	4	2.50	2310	5813	2310	0	2014
21610	402101.0	2	0.75	1020	1350	1020	0	2009
21611	400000.0	3	2.50	1600	2388	1600	0	2004
21612	325000.0	2	0.75	1020	1076	1020	0	2008

21613 rows × 8 columns

PRACTICE OPPORTUNITY #2 [OPTIONAL]:

- Using Seaborn, plot the pairplot for the features contained in "house_df_sample"
- Explore the data and perform sanity check

```
In [ ]:
```

TASK #4: PERFORM DATA CLEANING AND FEATURE ENGINEERING

```
In [12]: selected_features = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'sqft_above', 'sqft_basement']
```

```
In [13]: X = house_df[selected_features]
```

```
In [14]: X
```

```
Out[14]:
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_above	sqft_basement
0	3	1.00	1180	5650	1.0	1180	0
1	3	2.25	2570	7242	2.0	2170	400
2	2	1.00	770	10000	1.0	770	0
3	4	3.00	1960	5000	1.0	1050	910
4	3	2.00	1680	8080	1.0	1680	0
...
21608	3	2.50	1530	1131	3.0	1530	0
21609	4	2.50	2310	5813	2.0	2310	0
21610	2	0.75	1020	1350	2.0	1020	0
21611	3	2.50	1600	2388	2.0	1600	0
21612	2	0.75	1020	1076	2.0	1020	0

21613 rows × 7 columns

```
In [15]: y = house_df['price']
```

```
In [16]: y
```

```
Out[16]: 0      221900.0  
1      538000.0
```

```
In [17]: X.shape
```

```
Out[17]: (21613, 7)
```

```
In [18]: y.shape
```

```
Out[18]: (21613,)
```

```
In [19]: from sklearn.preprocessing import MinMaxScaler  
scaler = MinMaxScaler()  
X_scaled = scaler.fit_transform(X)
```

```
In [20]: X_scaled
```

```
Out[20]: array([[0.09090909, 0.125      , 0.06716981, ..., 0.      , 0.09758772,  
0.        ],  
[0.09090909, 0.28125    , 0.17207547, ..., 0.4      , 0.20614035,  
0.08298755],  
[0.06060606, 0.125      , 0.03622642, ..., 0.      , 0.05263158,  
0.        ],  
...,  
[0.06060606, 0.09375    , 0.05509434, ..., 0.4      , 0.08004386,  
0.        ],  
[0.09090909, 0.3125    , 0.09886792, ..., 0.4      , 0.14364035,  
0.        ],  
[0.06060606, 0.09375    , 0.05509434, ..., 0.4      , 0.08004386,  
0.        ]])
```

```
In [21]: X_scaled.shape
```

```
Out[21]: (21613, 7)
```

```
In [22]: scaler.data_max_
```

```
Out[22]: array([3.300000e+01, 8.000000e+00, 1.354000e+04, 1.651359e+06,  
3.500000e+00, 9.410000e+03, 4.820000e+03])
```

```
In [23]: scaler.data_min_
```

```
Out[23]: array([ 0.,  0., 290., 520.,  1., 290.,  0.])
```

```
In [24]: y = y.values.reshape(-1,1)
```

```
In [25]: y_scaled = scaler.fit_transform(y)
```

```
In [26]: y_scaled
```

```
Out[26]: array([[0.01926557],  
[0.06072131],  
[0.01377049],  
...,  
[0.04289849],  
[0.04262295],  
[0.03278689]])
```

TASK #5: TRAIN A DEEP LEARNING MODEL WITH LIMITED NUMBER OF FEATURES

```
In [27]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_scaled, test_size = 0.25)
```

```
In [28]: X_train.shape
```

```
Out[28]: (16209, 7)
```

```
In [29]: X_test.shape
```

```
Out[29]: (5404, 7)
```

```
In [30]: import tensorflow.keras  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense  
  
model = Sequential()  
model.add(Dense(100, input_dim = 7, activation = 'relu'))  
model.add(Dense(100, activation='relu'))  
model.add(Dense(100, activation='relu'))  
model.add(Dense(1, activation = 'linear'))
```

```
In [31]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 100)	800
dense_1 (Dense)	(None, 100)	10100
dense_2 (Dense)	(None, 100)	10100
dense_3 (Dense)	(None, 1)	101
=====		

Total params: 21,101

Trainable params: 21,101

Non-trainable params: 0

```
In [32]: model.compile(optimizer = 'Adam', loss = 'mean_squared_error')
```

```
In [33]: epochs_hist = model.fit(X_train, y_train, epochs = 100, batch_size = 50, validation_split = 0.2)
```

Train on 12967 samples, validate on 3242 samples

Epoch 1/100
12967/12967 [=====] - 5s 374us/sample - loss: 0.0012 - val_loss: 0.0011
Epoch 2/100
12967/12967 [=====] - 1s 64us/sample - loss: 0.0011 - val_loss: 0.0010
Epoch 3/100
12967/12967 [=====] - 1s 56us/sample - loss: 0.0011 - val_loss: 0.0012
Epoch 4/100
12967/12967 [=====] - 1s 56us/sample - loss: 0.0010 - val_loss: 0.0010
Epoch 5/100
12967/12967 [=====] - 1s 65us/sample - loss: 0.0010 - val_loss: 0.0010
Epoch 6/100
12967/12967 [=====] - 1s 62us/sample - loss: 0.0010 - val_loss: 0.0010
Epoch 7/100
12967/12967 [=====] - 1s 60us/sample - loss: 0.0010 - val_loss: 0.0010
Epoch 8/100
12967/12967 [=====] - 1s 56us/sample - loss: 0.0010 - val_loss: 0.0011
Epoch 9/100
12967/12967 [=====] - 1s 72us/sample - loss: 9.9609e-04 - val_loss: 0.0010
Epoch 10/100
12967/12967 [=====] - 1s 65us/sample - loss: 9.7768e-04 - val_loss: 0.0011
Epoch 11/100
12967/12967 [=====] - 1s 60us/sample - loss: 9.8589e-04 - val_loss: 0.0011
Epoch 12/100
12967/12967 [=====] - 1s 55us/sample - loss: 9.8197e-04 - val_loss: 0.0010
Epoch 13/100
12967/12967 [=====] - 1s 59us/sample - loss: 9.7691e-04 - val_loss: 0.0010
Epoch 14/100
12967/12967 [=====] - 1s 70us/sample - loss: 9.9276e-04 - val_loss: 0.0011
Epoch 15/100
12967/12967 [=====] - 1s 61us/sample - loss: 9.7304e-04 - val_loss: 0.0010
Epoch 16/100
12967/12967 [=====] - 1s 64us/sample - loss: 9.6194e-04 - val_loss: 0.0010
Epoch 17/100
12967/12967 [=====] - 1s 58us/sample - loss: 9.5018e-04 - val_loss: 0.0011
Epoch 18/100
12967/12967 [=====] - 1s 62us/sample - loss: 9.4871e-04 - val_loss: 0.0010
Epoch 19/100
12967/12967 [=====] - 1s 59us/sample - loss: 9.5959e-04 - val_loss: 0.0011
Epoch 20/100
12967/12967 [=====] - 1s 60us/sample - loss: 9.5551e-04 - val_loss: 0.0011
Epoch 21/100
12967/12967 [=====] - 1s 60us/sample - loss: 9.6346e-04 - val_loss: 0.0010
Epoch 22/100
12967/12967 [=====] - 1s 64us/sample - loss: 9.4101e-04 - val_loss: 0.0010
Epoch 23/100
12967/12967 [=====] - 1s 57us/sample - loss: 9.5983e-04 - val_loss: 0.0010
Epoch 24/100
12967/12967 [=====] - 1s 64us/sample - loss: 9.5335e-04 - val_loss: 9.9948e-04
Epoch 25/100
12967/12967 [=====] - 1s 59us/sample - loss: 9.6045e-04 - val_loss: 1.0000e-03
Epoch 26/100
12967/12967 [=====] - 1s 57us/sample - loss: 9.2578e-04 - val_loss: 0.0010
Epoch 27/100
12967/12967 [=====] - 1s 57us/sample - loss: 9.6472e-04 - val_loss: 0.0010
Epoch 28/100
12967/12967 [=====] - 1s 57us/sample - loss: 9.4331e-04 - val_loss: 0.0011
Epoch 29/100
12967/12967 [=====] - 1s 60us/sample - loss: 9.3168e-04 - val_loss: 0.0010
Epoch 30/100
12967/12967 [=====] - 1s 60us/sample - loss: 9.6121e-04 - val_loss: 0.0011
Epoch 31/100
12967/12967 [=====] - 1s 87us/sample - loss: 9.3711e-04 - val_loss: 0.0010
Epoch 32/100
12967/12967 [=====] - 1s 106us/sample - loss: 9.4819e-04 - val_loss: 0.0011
Epoch 33/100
12967/12967 [=====] - 1s 70us/sample - loss: 9.2509e-04 - val_loss: 0.0010
Epoch 34/100
12967/12967 [=====] - 1s 62us/sample - loss: 9.2075e-04 - val_loss: 0.0010
Epoch 35/100
12967/12967 [=====] - 1s 61us/sample - loss: 9.2888e-04 - val_loss: 0.0010

TASK #6: EVALUATE TRAINED DEEP LEARNING MODEL PERFORMANCE

```
In [34]: epochs_hist.history.keys()
```

```
Out[34]: dict_keys(['loss', 'val_loss'])
```

```
In [35]: plt.plot(epochs_hist.history['loss'])
plt.plot(epochs_hist.history['val_loss'])
plt.title('Model Loss Progress During Training')
plt.xlabel('Epoch')
plt.ylabel('Training and Validation Loss')
plt.legend(['Training Loss', 'Validation Loss'])
```

```
Out[35]: <matplotlib.legend.Legend at 0x20f67421358>
```



```
In [36]: # 'bedrooms','bathrooms','sqft_living','sqft_lot','floors', 'sqft_above', 'sqft_basement'
X_test_1 = np.array([[ 4,  3, 1960, 5000, 1, 2000, 3000 ]])
```

```
scaler_1 = MinMaxScaler()
X_test_scaled_1 = scaler_1.fit_transform(X_test_1)
```

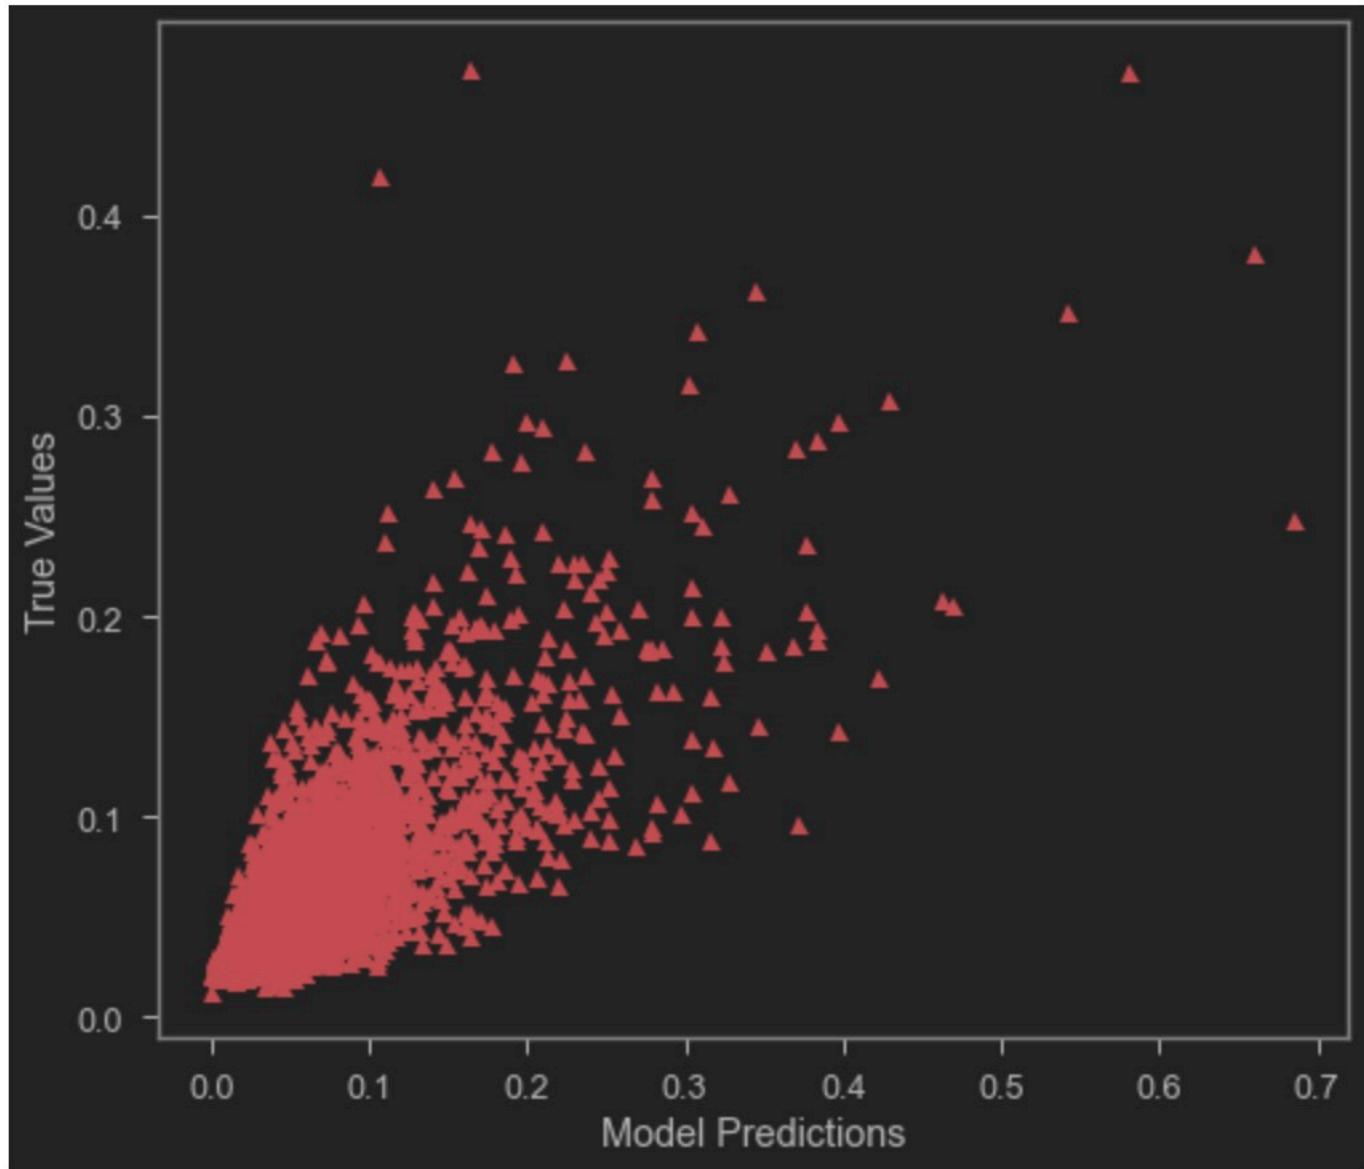
```
y_predict_1 = model.predict(X_test_scaled_1)
```

```
y_predict_1 = scaler.inverse_transform(y_predict_1)
y_predict_1
```

```
Out[36]: array([[218712.4]], dtype=float32)
```

```
In [37]: y_predict = model.predict(X_test)
plt.plot(y_test, y_predict, "^^", color = 'r')
plt.xlabel('Model Predictions')
plt.ylabel('True Values')
```

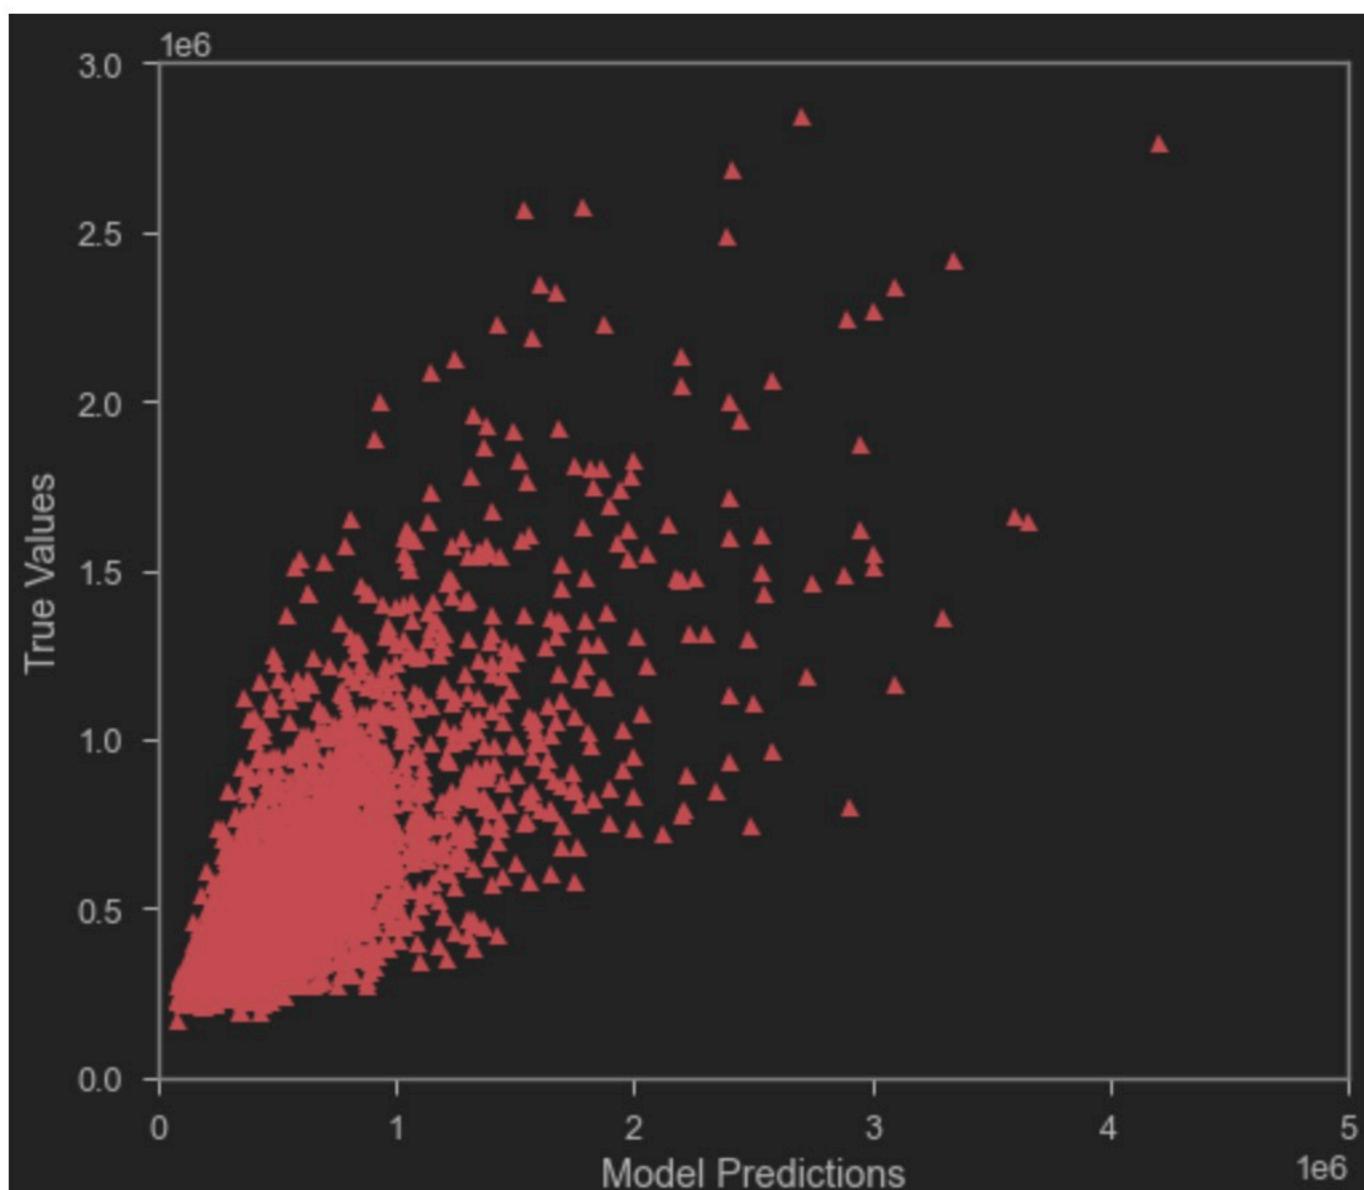
```
Out[37]: Text(0, 0.5, 'True Values')
```



```
In [38]: y_predict_orig = scaler.inverse_transform(y_predict)
y_test_orig = scaler.inverse_transform(y_test)
```

```
In [39]: plt.plot(y_test_orig, y_predict_orig, "^^", color = 'r')
plt.xlabel('Model Predictions')
plt.ylabel('True Values')
plt.xlim(0, 5000000)
plt.ylim(0, 3000000)
```

```
Out[39]: (0.0, 3000000.0)
```



```
In [40]: k = X_test.shape[1]
n = len(X_test)
n
```

```
Out[40]: 5404
```

```
In [41]: k
```

```
Out[41]: 7
```

```
In [42]:
```

```
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
from math import sqrt

RMSE = float(format(np.sqrt(mean_squared_error(y_test_orig, y_predict_orig)), '.3f'))
MSE = mean_squared_error(y_test_orig, y_predict_orig)
MAE = mean_absolute_error(y_test_orig, y_predict_orig)
r2 = r2_score(y_test_orig, y_predict_orig)
adj_r2 = 1-(1-r2)*(n-1)/(n-k-1)

print('RMSE =', RMSE, '\nMSE =', MSE, '\nMAE =', MAE, '\nR2 =', r2, '\nAdjusted R2 =', adj_r2)
```

RMSE = 244520.289

MSE = 59790171709.68345

MAE = 153968.88459937082

R2 = 0.5680483772792423

Adjusted R2 = 0.5674880249147046

TASK #7. TRAIN AND EVALUATE A DEEP LEARNING MODEL WITH INCREASED NUMBER OF FEATURES (INDEPENDANT VARIABLES)

```
In [68]: selected_features = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'sqft_above', 'sqft_basement', 'water_yr_renovated', 'zipcode', 'lat', 'long', 'sqft_living15', 'sqft_lot15']

X = house_df[selected_features]
```

```
In [69]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [70]: y = house_df['price']
```

```
In [71]: y = y.values.reshape(-1,1)
y_scaled = scaler.fit_transform(y)
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_scaled, test_size = 0.25)
```

```
In [72]: import tensorflow.keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

model = Sequential()
model.add(Dense(10, input_dim = 19, activation = 'relu'))
model.add(Dense(10, activation = 'relu'))
model.add(Dense(1, activation = 'linear'))
```

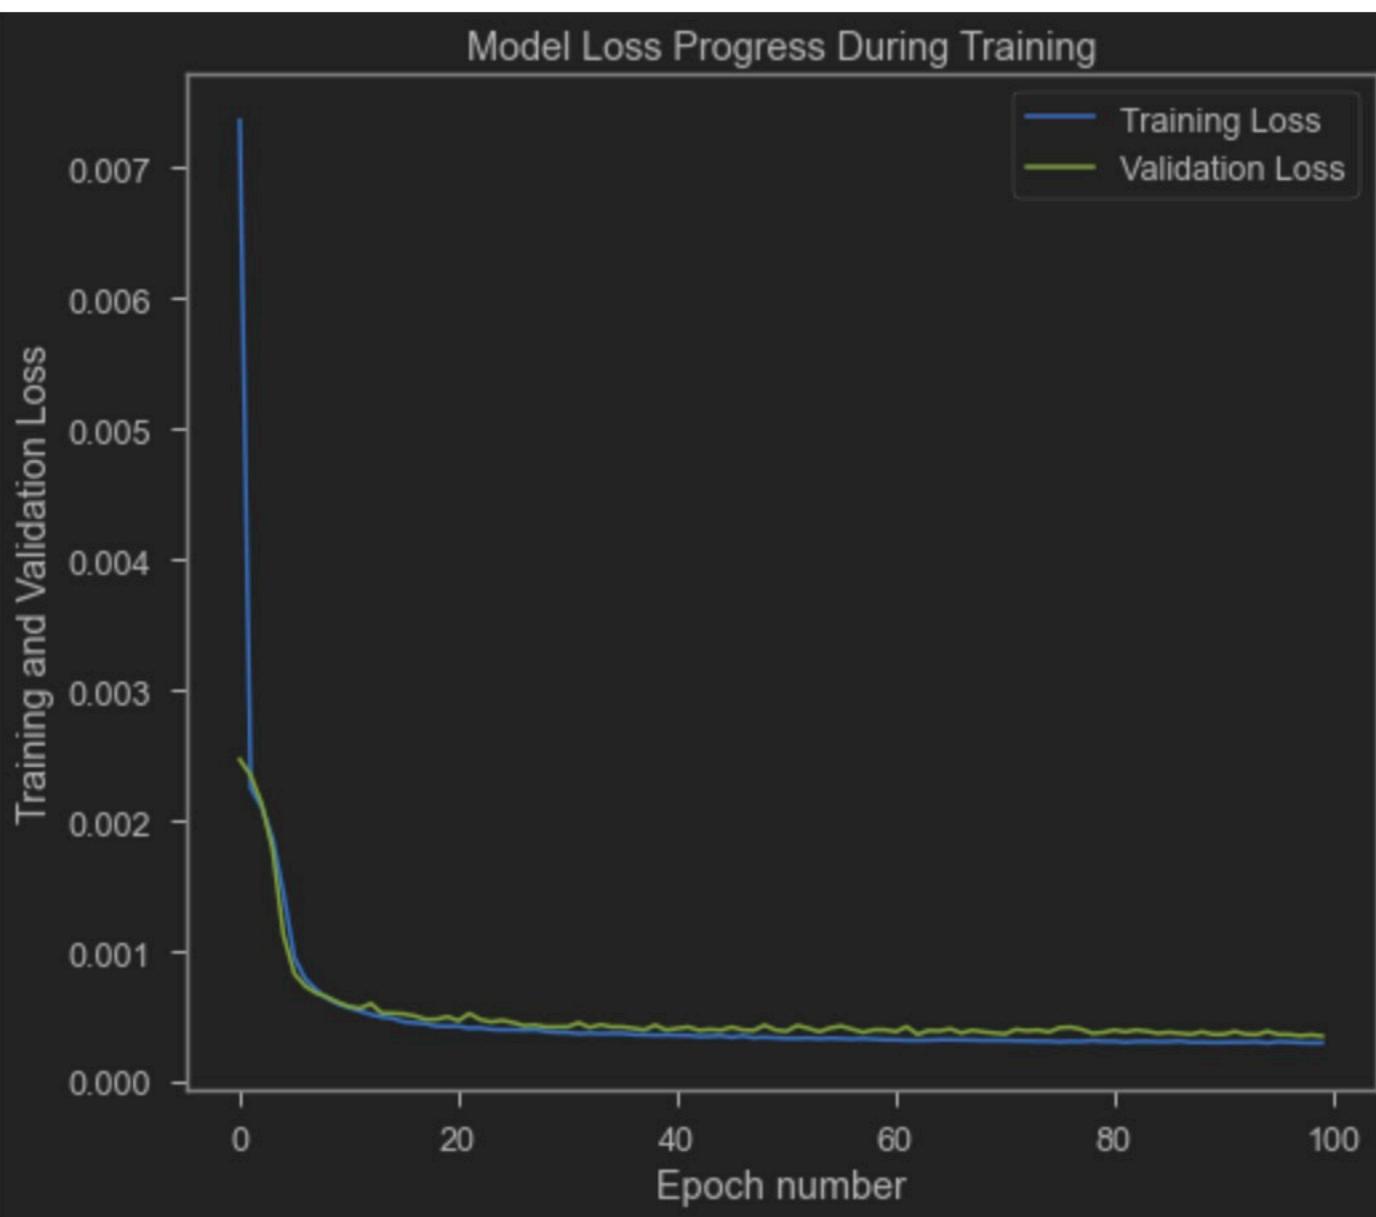
```
In [73]: model.compile(optimizer = 'adam', loss = 'mean_squared_error')
```

```
In [74]: epochs_hist = model.fit(X_train, y_train, epochs = 100, batch_size = 50, verbose = 1, validation_split = 0.2)
```

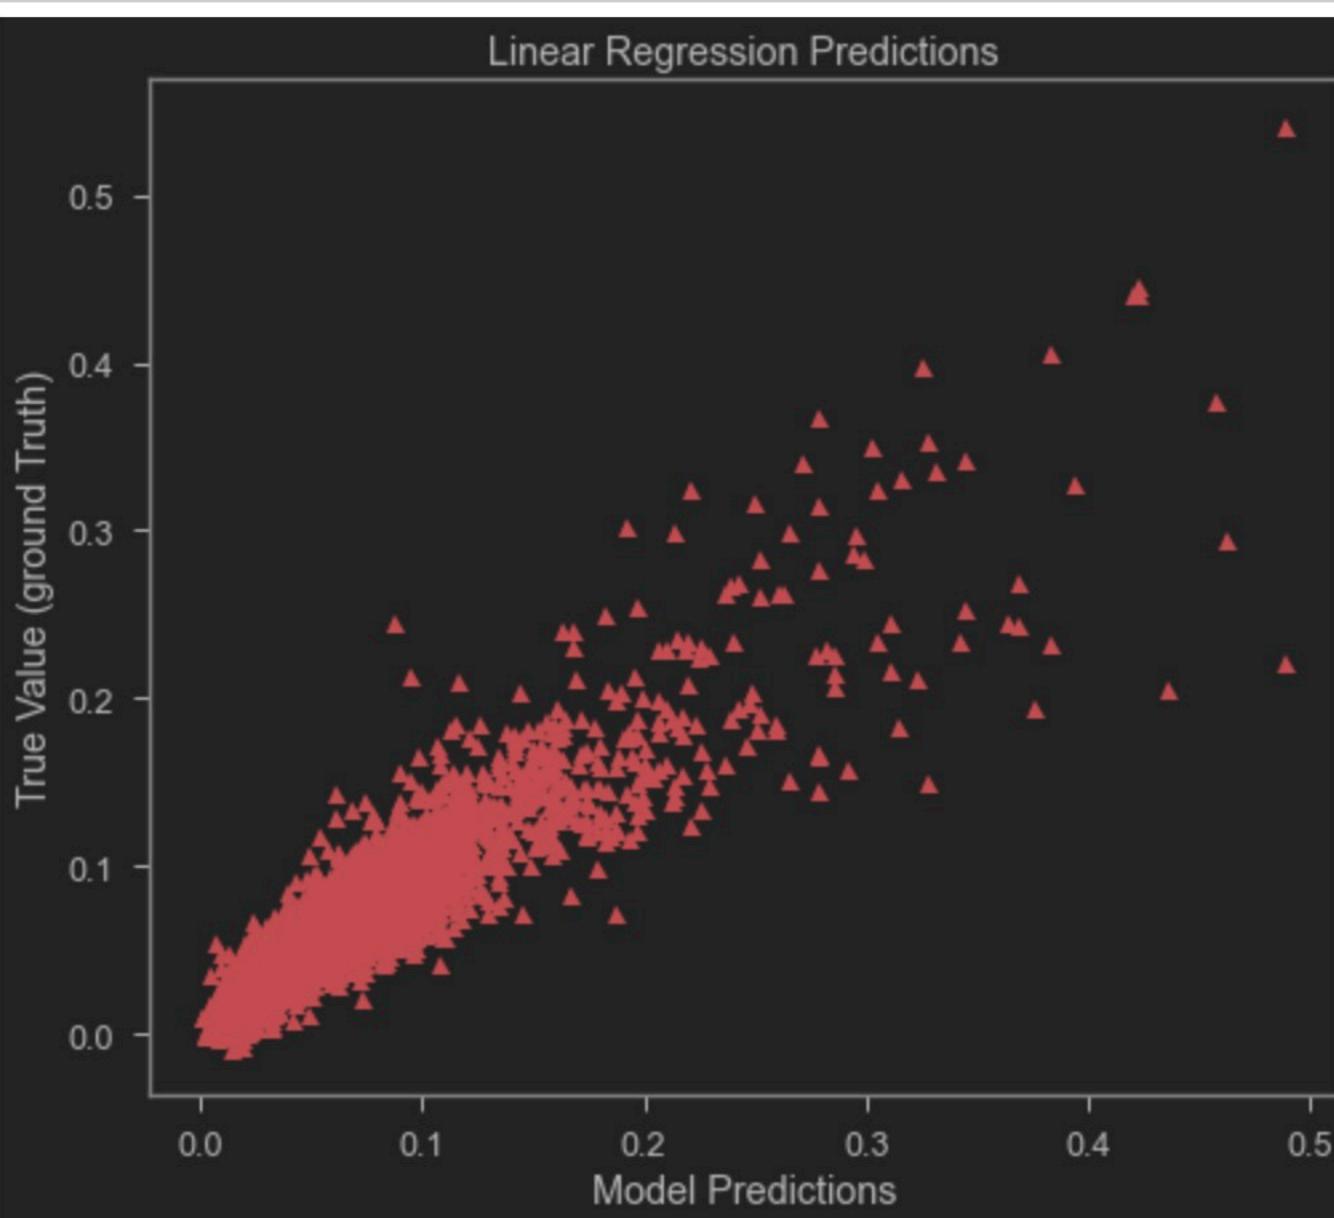
```
Train on 12967 samples, validate on 3242 samples
Epoch 1/100
12967/12967 [=====] - 2s 119us/sample - loss: 0.0074 - val_loss: 0.0025
Epoch 2/100
12967/12967 [=====] - 1s 61us/sample - loss: 0.0022 - val_loss: 0.0023
Epoch 3/100
12967/12967 [=====] - 1s 52us/sample - loss: 0.0021 - val_loss: 0.0021
Epoch 4/100
12967/12967 [=====] - 1s 52us/sample - loss: 0.0019 - val_loss: 0.0018
Epoch 5/100
12967/12967 [=====] - 1s 55us/sample - loss: 0.0014 - val_loss: 0.0011
Epoch 6/100
12967/12967 [=====] - 1s 62us/sample - loss: 9.4983e-04 - val_loss: 8.2648e-04
Epoch 7/100
12967/12967 [=====] - 1s 68us/sample - loss: 7.8759e-04 - val_loss: 7.3160e-04
Epoch 8/100
12967/12967 [=====] - 1s 63us/sample - loss: 7.0361e-04 - val_loss: 6.7910e-04
Epoch 9/100
12967/12967 [=====] - 1s 68us/sample - loss: 6.3749e-04 - val_loss: 6.4892e-04
Epoch 10/100
```

```
In [75]: plt.plot(epochs_hist.history['loss'])
plt.plot(epochs_hist.history['val_loss'])
plt.title('Model Loss Progress During Training')
plt.ylabel('Training and Validation Loss')
plt.xlabel('Epoch number')
plt.legend(['Training Loss', 'Validation Loss'])
```

```
Out[75]: <matplotlib.legend.Legend at 0x20f67945748>
```



```
In [79]: y_predict = model.predict(X_test)
plt.plot(y_test, y_predict, '^', color = 'r')
plt.xlabel("Model Predictions")
plt.ylabel("True Value (ground Truth)")
plt.title('Linear Regression Predictions')
plt.show()
```



```
In [80]: y_predict_orig = scaler.inverse_transform(y_predict)
y_test_orig = scaler.inverse_transform(y_test)
```

```
In [81]: from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
from math import sqrt

RMSE = float(format(np.sqrt(mean_squared_error(y_test_orig, y_predict_orig)), '.3f'))
MSE = mean_squared_error(y_test_orig, y_predict_orig)
MAE = mean_absolute_error(y_test_orig, y_predict_orig)
r2 = r2_score(y_test_orig, y_predict_orig)
adj_r2 = 1-(1-r2)*(n-1)/(n-k-1)

print('RMSE =', RMSE, '\nMSE =', MSE, '\nMAE =', MAE, '\nR2 =', r2, '\nAdjusted R2 =', adj_r2)
```

```
RMSE = 139326.953
MSE = 19411999744.979095
MAE = 85971.19128155938
R2 = 0.8420021760623091
Adjusted R2 = 0.8417972122432646
```

PRACTICE OPPORTUNITY #4 [OPTIONAL]:

- Change the architecture of the network to increase the coefficient of determination to at least 0.86.

```
In [ ]:
```

GREAT JOB!